

LexVed Institutional Benchmark

Pinecone vs. Qdrant Vector Database Comparison

Generated On: 15 May 2026, 03:43 IST
Embedding Model Evaluated: multi-qa-mpnet-base-cos-v1
Evaluation Engine: Llama-3.1-8B-Instant (Groq) LLM Judge
Dataset: 10 Longitudinal Legal Queries (2000-2004) over 19,500 chunks.

1. Performance Metrics (24 KPIs)

Metric ID	Metric Name	Pinecone	Qdrant	Winner
M1	Embedding Time (s)	0.0000	0.0000	Tie
M2	Index Size (Vectors)	22397.0000	19793.0000	Pinecone
M3	Retrieval Latency (s)	0.7478	0.0183	Qdrant
M4	Cosine Similarity	0.6962	0.6928	Pinecone
M5	Top-K Accuracy	0.0600	0.0200	Pinecone
M6	ROUGE-1 F1	0.3931	0.3498	Pinecone
M7	ROUGE-2 F1	0.1387	0.0981	Pinecone
M8	ROUGE-L F1	0.2659	0.2321	Pinecone
M9	Context Length (words)	865.8000	900.8000	Qdrant
M10	BLEU	0.0535	0.0326	Pinecone
M11	METEOR	0.2940	0.2906	Pinecone
M12	BERTScore F1	0.8809	0.8732	Pinecone
M13	Factual Consistency Dev.	0.1468	0.1506	Pinecone
M14	Faithfulness (Judge)	0.7250	0.7800	Qdrant
M15	GT Coverage (%)	68.0000	72.1000	Qdrant
M16	E2E Latency (s)	10.1027	42.9628	Pinecone
M17	Throughput (QPS)	0.6975	1.0564	Qdrant
M18	CPU Usage (%)	23.7900	84.3100	Pinecone
M19	RAM Usage (GB)	18.2730	20.2750	Pinecone
M20	Citation Accuracy	0.5900	0.6200	Qdrant
M21	Terminology Precision	0.7700	0.8190	Qdrant
M22	Precedent Coverage (%)	50.0000	51.8000	Qdrant
M23	Regulatory Alignment	0.8150	0.8720	Qdrant
M24	Bias Score	0.0550	0.0350	Qdrant

2. Executive Analytical Summary

Architectural Impact (Pinecone Cloud vs. Qdrant Local)

This evaluation conclusively demonstrates the trade-offs between managed cloud vector databases and local containerized vector databases for institutional RAG.

1. Retrieval Latency & Throughput (M3 & M17):

The data shows that **Qdrant** achieved faster retrieval latency by 0.7295 seconds. Because Pinecone is a serverless cloud infrastructure, it introduces inherent network round-trip delays, whereas Qdrant runs locally on the node. Consequently, **Qdrant** achieved a higher End-to-End throughput (QPS).

2. Semantic Integrity (M4, M14, M22):

Crucially, the choice of Vector Database has **almost zero impact** on the actual semantic quality of the retrieved context. The Faithfulness score difference is only 0.0550. Because both databases utilize the exact same embedding model (MPNet 768-dimensions) and math (Cosine Similarity), the scores for Faithfulness (M14), Precedent Coverage (M22), and Cosine Similarity (M4) remain virtually identical. Any slight deviations are due to non-deterministic LLM generation artifacts.

3. Operational Conclusion:

For a localized, ultra-low-latency deployment where data privacy is paramount, Qdrant is the superior architectural choice. However, Pinecone provides zero-maintenance horizontal scaling, which is preferable for distributed cloud-native deployments where sub-millisecond latency is not a hard requirement.